

CANOPY. - Product Brief

Prime Three Offerings (v1)

CANOPY builds systems that model real-world environments as they evolve, using deployable sensing to create continuously updated synthetic models that inform real-time decisions.

1 Synthetic Models (Core Platform)

A software platform that maintains living, time-aware models of physical environments by fusing live sensor data, UAV data, and external data sources. This is not a static digital twin. It is a continuously updating synthetic representation of reality.

Inputs

- UAV data (imagery, LiDAR, telemetry)
- Ground / temporary sensor nodes (environmental, vibration, etc.)
- External APIs (weather, mapping, historical datasets)
- Manual annotations (optional)

Core capabilities (v1)

- Real-time state representation (what is happening now)
- Time-series persistence (what has happened)
- Predictive modelling hooks (what may happen next)
- Spatial awareness (GIS-native)
- Uncertainty scoring (confidence per data layer)

Outputs

- Model state API
- Time-based simulations
- Data layers consumable by dashboards and external systems

Non-goals (v1)

- Photorealistic rendering
- Game-engine visuals
- AI-only “black box” predictions

Why it exists

This is the long-term value moat. Everything else feeds or consumes this layer.

2 Deployable Sensing (Truth Layer)

A set of tools and workflows for rapidly collecting ground-truth data in environments where permanent infrastructure does not exist.

This includes UAVs and temporary sensors, but the value is in deployment + ingestion, not the hardware itself.

Components

- UAV-based data acquisition
- Temporary / mobile ground sensors
- Mesh or backhaul communications
- Location-aware data tagging

Core capabilities (v1)

- Rapid deployment (minutes, not weeks)
- Automatic geospatial tagging
- Time-synchronised data streams
- Health/status reporting for sensors
- Seamless ingestion into Synthetic Models

Outputs

- Live sensor streams
- Buffered data uploads
- Sensor metadata (location, uptime, confidence)

Non-goals (v1)

- Mass-manufactured sensor hardware
- Long-term fixed infrastructure replacement
- Perfect accuracy (confidence is surfaced instead)

Why it exists

Synthetic Models are only as good as the data feeding them.
This layer ensures CANOPY controls data quality and relevance.

3 Operational Intelligence (Decision Layer)

What it is

The interface layer that turns Synthetic Models into actionable understanding for operators, planners, and decision-makers.

This is not an analytics tool for data scientists. It is built for people who need to act.

Core capabilities (v1)

- Live situational awareness views
- Alerts based on model thresholds
- Time-based playback (“what changed”)
- Scenario comparison (baseline vs current)
- Integration with existing GIS / ops tools

Outputs

- Dashboards
- Alerts / notifications
- Reports
- External system integrations

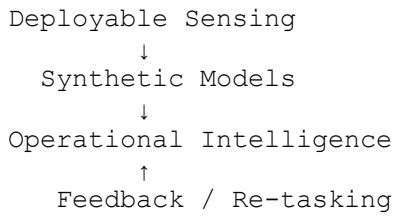
Non-goals (v1)

- Custom dashboards for every user
- Heavy workflow engines
- Complex permissions hierarchies

Why it exists

This is how CANOPY becomes used daily, not just deployed occasionally.

System Relationship (Critical)



- Sensing feeds models
- Models inform decisions
- Decisions drive new sensing

This feedback loop is the product.

Technical Principles (Guiding Rules)

- API-first architecture
- GIS-native data structures
- Time-aware by default
- Confidence and uncertainty are first-class data
- Modular (each layer usable independently)
- Dark-mode, low-distraction UI

Focus Summary

Build this first

- Ingestion → Model → Visualisation loop
- One environment
- One sensing modality
- One predictive output

Explicitly defer

- Scale optimisation
- Full automation
- Advanced AI/ML
- Broad hardware support

Model reality accurately, update it continuously, and make it usable.